

# Look Before You Leap: A Benchmark and Exploratory Study of Uncertainty Analysis for Large Language Models

Yuheng Huang, Jiayang Song, Zhijie Wang, Shengming Zhao,  
Huaming Chen, Felix Juefei-Xu, Lei Ma ✉

**Abstract**—The recent performance leap of Large Language Models (LLMs) opens up new opportunities across numerous industrial applications and domains. However, the potential erroneous behavior (e.g., the generation of misinformation and hallucination) has also raised severe concerns for the trustworthiness of LLMs, especially in safety-, security- and reliability-sensitive industrial scenarios, potentially hindering real-world adoptions. While uncertainty estimation has shown its potential for interpreting the prediction risks made by classic machine learning (ML) models, the unique characteristics of recent LLMs (e.g., adopting self-attention mechanism as its core, very large-scale model size, often used in generative contexts) pose new challenges for the behavior analysis of LLMs. Up to the present, little progress has been made to better understand whether and to what extent uncertainty estimation can help characterize the capability boundary of an LLM, to counteract its undesired behavior, which is considered to be of great importance with the potential wide-range applications of LLMs across industry domains. To bridge the gap, in this paper, we initiate an early exploratory study of the risk assessment of LLMs from the lens of uncertainty. In particular, we conduct a large-scale experimental study with as many as twelve uncertainty estimation methods and five general LLMs on four NLP tasks and five code-specific LLMs on two code generation tasks to investigate to what extent uncertainty estimation techniques could help characterize the prediction risks of LLMs. Our findings confirm the potential of uncertainty estimation for revealing LLMs’ uncertain/non-factual predictions. The insights derived from our study can pave the way for more advanced analysis and research on LLMs, ultimately aiming at enhancing their trustworthiness.

**Index Terms**—Large Language Models, Deep Neural Networks, Uncertainty Estimation, Software Reliability

## 1 INTRODUCTION

Large Language Models (LLMs) have demonstrated impressive capabilities in miscellaneous Natural Language Processing (NLP) tasks and promising adaptability in practical applications across diverse domains, including but not limited to content moderation [1], code generation [2], conversational AI [3], and personalized content recommendations [4]. The scale of deployment is vast, addressing the needs of diverse user demographics and industries. Among the leading AI corporations, our industrial partner Meta has consistently championed the progression of pertinent technology, launching foundational models such as Llama [5] and Llama-2 [6]. As of now, there are more than 3,500 enterprise project initiations and over 7,000 GitHub projects rooted in the Llama series models [7].

Despite the attractive performance that LLMs present and their rapid evolution within both academics and industries, an urgent common concern about LLMs has the propensity of generating erroneous information without warning. Such phenomenon of erroneous generation can

exhibit in terms of different manifestations (e.g., hallucination [8], disinformation [9], bias [10]) across various tasks. In general, the current LLMs are found to have the tendency to generate problematic, nonfactual responses that are not from training sources or misguided by biased data. However, these responses are often presented in a natural human-like tone [11], [12]. Such characteristics cause erroneous information to be highly mixed and intertwined with confident and factual contexts, making their detection and localization difficult without close inspection and diligent fact-checking [13]. As an example, Fig. 2 (GPT-3) depicts an example of an LLM answering a question with nonfactual information.

Risk assessments thus become crucial in the process of mitigating such threats. A recent survey highlights that 98% of respondents, encompassing domain experts and civil society members, firmly believe that AGI (artificial general intelligence) labs should undertake risk assessments before deployment [14]. For AI industry, implementing comprehensive risk assessment methods is not just a technical necessity but also an ethical obligation. Major tech corporations [15] such as Microsoft [16], OpenAI [17], Amazon [18], and Google [19], along with non-governmental organizations (NGOs, e.g., the Centre for the Governance of AI [20]), are fervently working towards developing safe, secure, transparent, reliable and responsible LLMs and AGI applications. As a driving force behind open and collaborative AI research, Meta also commits substantial resources

✉ Corresponding author

• Yuheng Huang, Jiayang Song, Zhijie Wang and Shengming Zhao are with the University of Alberta, Canada. E-mail: {yuheng18, jiaayan13, zhijie.wang, shengmi1}@ualberta.ca

• Huaming Chen is with the University of Sydney, Australia. E-mail: huaming.chen@sydney.edu.au

• Felix Juefei-Xu is with GenAI at Meta, USA. E-mail: felixxu@meta.com

• Lei Ma is with The University of Tokyo, Japan, and University of Alberta, Canada. E-mail: ma.lei@acm.org

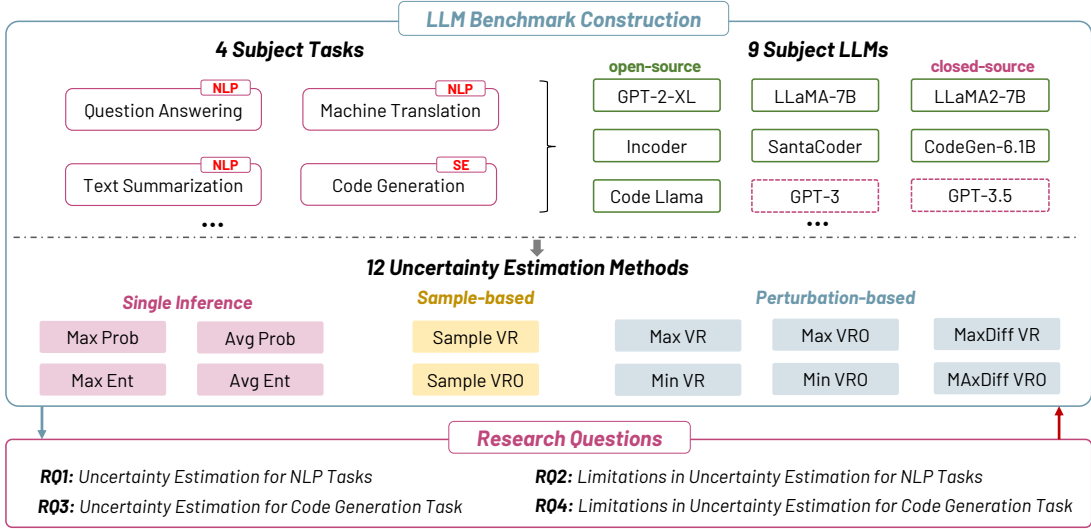


Fig. 1: A high-level overview and workflow of this paper.

**Example:** Question Answering with correct response and risky response




Fig. 2: Uncertainty estimation for a QA task.

to the development of responsible AI, emphasizing trustworthiness, transparency, robustness, etc. [21], [22]. These endeavors encompass a range of reports (e.g., *Building Generative AI Responsibly* [23]), open-source tools (e.g., the model interpretability framework *Captum* [24]), and datasets (e.g., *Hateful Memes* [25]).

Although there has been substantial work on other AI models, risk assessments for LLMs are still in their infancy. Due to their billions of parameters, vast amounts of (often inaccessible) data, and potential closed-source nature, LLMs present significant challenges for analysis and safeguard. Yet, such a safeguard is crucial, especially considering the widespread adoption of LLMs. Stakeholders from research, industry, open source initiatives, NGOs, and businesses may all be negatively affected by untrustworthy LLMs. Uncertainty estimation, aimed at gauging the confidence level of model outputs [26], [27], [28], stands out as a promising approach for identifying risks in general Machine Learning (ML) models. Such techniques also have the potential for detecting erroneous generation from LLMs [29] even under black-box settings. It is thus possible to take them as plug-and-play tools in both academic and industrial scenarios. For example, Fig. 2 shows that a higher uncertainty score could possibly indicate an erroneous generation of an LLM. However, it is still unclear whether and to what extent uncertainty estimation methods could do when measuring and characterizing an LLM’s capability limitations. Furthermore, it also raises questions such as “Are there better practices for employing these methods in practical scenarios?” , “Do we

need further adaptations from the industrial perspective to better cater to the distinct features of LLMs (e.g., task diversity and high computational cost)?”, etc. To the best of our knowledge, up to the present, there is a lack of a general framework that integrates different uncertainty estimation methods for LLMs, as well as a systematic study to investigate the effectiveness of uncertainty estimation in characterizing an LLMs’ capabilities.

To bridge this gap, in this paper, we present an exploratory study to understand the trustworthiness of LLMs from the lens of uncertainty estimation. Considering the generality and versatility for various application scenarios, we strive to identify suitable methods to minimize the requirement of LLMs internal information (e.g., model architecture, model parameters). Such criteria enable these methods to be seamlessly adapted and incorporated by end-users of commercial models, such as GPT3.5. Overall, we collect and implement as many as 12 representative uncertainty estimation methods that were originally designed for general DNNs, and successfully adapt them to the contexts of LLM applications. To better capture an in-depth understanding of the effectiveness of these methods, we conduct large-scale experiments with as many as nine LLMs on both NLP (i.e., question answering, text summarization, machine translation) and software programming (i.e., code generation) tasks to analyze the correlation between uncertainty estimation results and LLMs performance. The models comprise four from *MetaAI*, three from *OpenAI*, one from *Salesforce*, and one from *BigCode*. The overall workflow of our work is shown in Fig. 1. In particular, we investigate the following research questions:

- **RQ1**: To what extent can the uncertainty estimation techniques help identify potential risks of LLMs in NLP tasks?
- **RQ2**: What limitations do the uncertainty estimation methods encounter when applied to LLMs in the context of NLP tasks?
- **RQ3**: To what extent can the uncertainty estimation methods assist in identifying potential risks of LLMs for code generation?
- **RQ4**: What potential limitation do the uncertainty estima-

tion methods face when being applied to LLMs for code generation?

Our findings validate that uncertainty measurement can, to an extent, be helpful in detecting erroneous responses in general NLP tasks. Additionally, it has also shown to be promising as an indicator for pinpointing faulty programs produced by LLMs. Even though, these methods might fall short in detecting nuanced errors made by high-performance commercial models. They seem better suited for filtering out more blatant mistakes. Based on the results, we further discuss the insights from our study and highlight a few potential research directions of leveraging uncertainty estimation to enhance the trustworthiness of LLMs for real-world applications across domains. *First*, research efforts are needed with novel uncertainty estimation techniques exclusively for LLMs to better fit the corresponding diverse task-handling ability. *Second*, we observe that different LLMs can sometimes display markedly distinct uncertain behaviors. Consequently, even though these methods are inherently model-agnostic, stakeholders may need to undertake model-specific optimizations to achieve enhanced performance. *Furthermore*, we observe that the prompt template used in the reinforcement learning from human feedback (RLHF [30]) could potentially impact the accuracy of uncertainty estimation.

The contributions of this paper are summarized as follows:

- We collected and implemented *twelve* different uncertainty estimation methods that are successfully adapted to enable the analysis of LLM, which are also applicable to both open-source and closed-source LLM models across different downstream tasks in the grey-box setting.
- We conducted a large-scale evaluation with nine LLMs on *six* tasks from *four* different domains.
- We provided an in-depth analysis of the challenges in existing uncertainty methods for LLMs and distilled a set of implications and future opportunities toward reliable and trustworthy LLMs.
- Our benchmark toolkit, encompassing the dataset, LLM inference, and uncertainty measurement protocols, will be made available for future research endeavors.

### The Contributions to the Software Engineering Field.

LLMs have revolutionized various aspects of software engineering [31], [32], including but not limited to automated code generation [33], [2], [34], [35], software testing [36], [37], [38], [39], debugging [40], program repair [41], [42], and document generation [43]. While LLMs can serve as a critical core for many new-era AI-enabled intelligent systems in the software engineering domain, their black-box nature and inherent uncertainties pose challenges for them to be applied in the real world in a transparent, reliable, safe, and secure way. It is thus urgent to investigate and explore effective quality assurance methods. Measuring uncertainty and taking it as an indicator of AI models' reliability has been studied extensively in the SE community [44], [45], [46]. While promising, most of them focus on classification tasks with relatively simple neural architectures. On the contrary, we initialize a very early stage study on autoregressive, large-scale language models and perform various uncertainty measurements across a wide spectrum of tasks.

We further provide more supplementary results and details as well as the source code to reproduce our study at our website: <https://sites.google.com/view/llm-uncertainty>.

## 2 BACKGROUND AND RELATED WORK

### 2.1 Large Language Models

In general, a *language model* models sequences of words as a probability distribution, which can be further used to generate coherent and contextually relevant text via conditioning based on a given prompt. Representative traditional language models include HMM (hidden Markov model) [48], *n*-gram [49], and RNN (recurrent neural networks) [50]. Recently, a specific type of neural network architecture, i.e., Transformer [51], has achieved attractive performance on language modelling. Large language models (LLMs) now typically refer to those Transformer-based language models pre-trained with large-scale text corpus and billions of parameters. LLMs have also achieved promising performance in many downstream tasks, e.g., text classification [52], text summarization [53], and machine translation [54].

Based on different Transformer architectures and pre-training tasks, LLMs largely fall into three categories: *encoder-only*, *encoder-decoder*, and *decoder-only*. *Encoder-only* LLMs also refer to masked language models (MLM), which are pre-trained through masking a certain portion of tokens (e.g., 15%) in a sequence. The training objective is to predict those masked tokens correctly. Representative *encoder-only* LLMs include BERT [55], RoBERTa [56], GraphCodeBERT [57], etc. Different from *encoder-only* LLMs, *encoder-decoder* LLMs, such as BART [58], CodeT5 [59], are pre-trained through masked span prediction (MSP). *Encoder-decoder* LLMs first learn a representation from input prompts before decoding into another sequence. They are thus usually trained for sequence-to-sequence purposes. Lately, *decoder-only* LLMs have become the mainstream of LLMs research due to their training efficiency and scalability for large-scale datasets and complex model architectures (i.e., billions of model parameters). *Decoder-only* LLMs are autoregressive models. Their training objective is to predict the next token given all previous (left-only) tokens. GPT-based (generative pre-trained transformers) models (e.g., GPT2 [60], GPT3 [47], LLaMA [5] and LLaMA-2 [6]) all belong to this category. In this work, we mainly focus on *decoder-only* LLMs since they have SOTA performance. We further detail subject LLMs in our study in Sec. 4.1.

Though LLMs are pre-trained without specific tasks in mind, they can often be used for downstream tasks in two ways, i.e., through (1) prompting and (2) fine-tuning. Prompting refers to the process of in-context learning that "teaches" an LLM to solve a specific task by injecting certain knowledge and instructions into the input prompts. Different from prompting, fine-tuning is the process of updating an LLM's neural network weights through a supervised learning strategy for certain tasks (e.g., text classification). Recently, *decoder-only* LLMs have also been used with reinforcement learning from human feedback (RLHF) to improve their performance in understanding complex input prompts and following human instructions. As a result, ChatGPT (GPT3.5 with RLHF) [3] has shown superior performance in solving various complex tasks (e.g., program

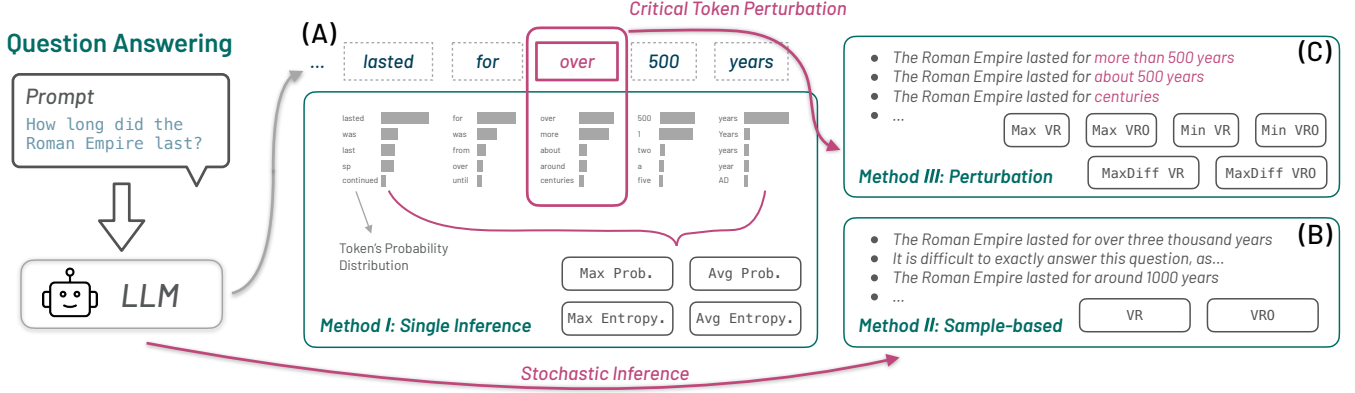


Fig. 3: A running example of how different uncertainty estimation methods work for a QA problem with GPT3 [47].

synthesizes [61], program repair [62]) by only following the user’s instructions and feedback through a dialogue. In our study, we mainly consider LLMs through prompting (that is widely used in practice) and focus on how to estimate an LLM’s uncertainty in a “black-box” way (i.e., only accessing the model’s prediction output probabilities). We introduce our experiment settings in Sec. 4.2.

## 2.2 Risk Assessment for ML Models

Machine learning (ML) models, especially deep learning ones, are known to be notoriously hard to interpret due to their complexity and opacity. Using ML models without appropriate risk assessment could potentially lead to particular concerns and threats regarding trustworthiness, e.g., safety [63], security [64], and ethics [65]. So far, some risk assessment techniques for general machine learning models have been proposed to reduce the impacts of these concerns [66], [67], [28], [68], [69]. Among them, there are two representative categories of risk assessment techniques: (1) data distribution analysis and (2) uncertainty estimation.

The data-driven nature of ML models requires developers to take data into account, especially when test data could be much more different compared with the training data in terms of their distribution. Data distribution analysis, including detecting distribution shift [67] and out-of-distribution samples [70], is proposed to identify such differences and avoid potential risks on unseen data. However, data distribution analysis usually requires access to training data, which is often not feasible for LLMs trained on either huge data corpus or private data corpus. Therefore, in this work, we propose to focus on *uncertainty estimation*.

In general, uncertainty estimation aims to measure an ML model’s confidence level of a certain prediction. There are two main types of uncertainty in an ML model’s predictions: *aleatoric* uncertainty and *epistemic* uncertainty [71]. *Aleatoric* uncertainty refers to the uncertainty that arises from observations (e.g., sensor noises in an ML model for autonomous driving). By contrast, *epistemic* uncertainty accounts for uncertainty in an ML model’s parameters. Insufficient knowledge of an ML model (e.g., lack of a specific type of training data) usually leads to high *epistemic* uncertainty. In this paper, we mainly discuss estimating *epistemic* uncertainty for LLMs. Uncertainty estimation roughly falls into

four categories [72]: (1) *single deterministic methods* [73], (2) *ensemble methods* [74], (3) *Bayesian methods* [75], and (4) *test-time augmentation methods* [76]. *Single deterministic methods* calculate prediction uncertainty based on one forward pass within a deterministic ML model. *Ensemble methods* estimate uncertainty based on a set of different ML models’ output. By contrast, *Bayesian methods* only leverage the inherent stochasticity of an ML model (e.g., dropout layer in deep neural networks [77]). *Test-time augmentation methods* are model-agnostic, which augment the input data at test-time to measure a model’s prediction uncertainty [78]. Since we focus on the risk assessment for one standalone LLM, *ensemble methods* are excluded from our study. We detail the uncertainty estimation methods used in this work in Sec. 3.

In addition to the aforementioned general risk assessment techniques, there are also a few works specified for risk assessment of LLMs [79], [80], [81], [82], [83], [84], [29]. The most related works are those proposed for uncertainty estimation of LLMs [82], [83], [84], [29]. Xiao et al. leverage *ensemble methods* to measure the natural language generation model’s uncertainty and detect potential hallucinations [82]. Similarly, Malinin et al. propose a unified uncertainty estimation method for autoregressive structured prediction tasks based on *ensemble methods* [83]. To overcome the challenge of capturing “semantic equivalence” in natural language, Kuhn et al. propose *semantic entropy* that incorporates linguistic invariances created by shared meanings [84]. Recently, Manakul et al. propose SelfCheckGPT, a black-Box hallucination detection method based on token-level prediction likelihood and entropy [29]. In light of the limitations of these works, our work is the first work that is not limited to specific natural language or tasks by covering twelve uncertainty estimation methods. Furthermore, our study investigates the role of uncertainty estimation with extensive experiments with nine LLMs and six tasks, providing insights and evidence for its effectiveness as the risk assessment technique for LLMs.

## 3 UNCERTAINTY ESTIMATION FOR LLMs

In this section, we first discuss the problem scenario in our study, including the corresponding assumptions. Then, we introduce our twelve uncertainty estimation techniques




Fig. 4: The available information we can get from the LLM in our setting. This is an illustration of how *LLAMA-7B* answers the question “when did the movie deep blue sea come out?”.

(three categories) based on the number of inferences required.

### 3.1 Problem Scenario

Given an input prompt  $X = [x_1, x_2, \dots, x_n]$  ( $x_i$  denotes  $i$ th input token), an LLM  $f$  with pre-trained weights  $w$  generates another sequence  $Y = [y_1, y_2, \dots, y_m]$  ( $y_j$  denotes  $j$ th generated token) through a decoding process:  $y_j = f([X, y_1, y_2, \dots, y_{j-1}]|w)$

An uncertainty estimation method  $g$  is to calculate a score  $u$  regarding the uncertainty of  $Y$ .

Though an LLM can be regarded as an ML model, it is limited by the inherent properties of some existing uncertainty estimation methods. Following we discuss the unique characteristics of LLMs and challenges in uncertainty estimation for LLMs compared with other ML models.

- **Complexity.** The state-of-the-art (SOTA) LLMs are usually pre-trained with billions of parameters (e.g., GPT-3 [47] model has 96 layers with 6.7 billion parameters). Therefore, “white-box” analysis for interpreting LLMs (e.g., inspecting neuron activation [85], inspecting attention values [86], [87]) that requires both significant manual and computational efforts is not feasible.
- **Opacity.** There is also a lack of opacity in SOTA LLMs. First, the SOTA LLMs are usually trained with large-scale text corpus, where such data can be either publicly available or from private sources. Therefore, risk assessment techniques that require access to training data (e.g., OOD detection) can not be used in our context. Additionally, some of the SOTA LLMs are potentially proprietary assets for a company (e.g., GPT-3 [47]), where one can only access the inference results through provided APIs.
- **Task diversity.** Though the usage of LLMs can be described in a general decoding form, tasks that LLMs can solve are of greater diversity. Notably, LLMs can be used for user-defined tasks through prompting/few-shot learning (Sec. 2.1). Therefore, uncertainty estimation methods that are proposed for a specific narrow domain (e.g., text classification) are hard to be used as a general risk assessment technique for LLMs.

With these characteristics and challenges in mind, we show the information that we can get from the majority of both open-source and closed-source LLMs in Fig. 3 (A) and explore possible existing solutions in this grey box setting for a general usage risk assessment. We summarize the

necessary information needed in this study in Fig. 4. In most cases, we are only able to obtain the output tokens as well as the probability distribution over the vocabulary for each token. Specifically, for closed-source models such as GPT-3 [47] and GPT-3.5 [88], one is only able to obtain the top- $k$  probabilities (i.e.,  $k$  possible tokens with the highest probabilities) for each predicted token. Formally, we can formulate the uncertainty estimation in our study as

$$u = g(f(\cdot), X, Y, P), \quad (1)$$

where  $P$  is either the probability distribution or top- $k$  probabilities for each token.

Finally, we select twelve uncertainty estimation methods covering *single deterministic methods*, *Bayesian methods*, and *test-time augmentation methods* (see discussion in Sec. 2.2). We categorize our uncertainty estimation methods based on the number of inferences required and detail them in the following.

### 3.2 Single-inference Uncertainty Estimation

Single-inference uncertainty estimation methods can be seen as *single deterministic methods* in Sec. 2.2. These methods usually calculate an ML model’s confidence based on the probability distribution of the prediction [89], [90]. In particular, such methods are usually used for classification tasks. Though LLMs can be used in various different tasks, the generation of each token can still refer to a classification problem (i.e., choose one token from the entire vocabulary). To aggregate the uncertainty information obtained at the token level, Manakul et al. [29] propose four different metrics to aggregate token-level uncertainty into sentence level.

In particular, a sentence-level uncertainty score can be obtained by taking either the maximum or average of the likelihood  $-\log p$  in a sentence:

$$\text{Max}(-\log p)_i = \max_j (-\log p_{ij}), \quad (2)$$

$$\text{Avg}(-\log p)_i = -\frac{1}{J} \sum_j \log p_{ij}, \quad (3)$$

where  $p_{ij}$  is the probability of a token at position  $j$  of a sentence  $i$ . Additionally, one can also replace the likelihood  $-\log p$  with the entropy  $\mathcal{H}$ :

$$\text{Max}(\mathcal{H})_i = \max_j [\mathcal{H}_{ij}], \quad (4)$$

$$\text{Avg}(\mathcal{H})_i = \frac{1}{J} \sum_j \mathcal{H}_{ij}, \quad (5)$$

where  $\mathcal{H}_{ij}$  is the entropy of this token’s probability distribution over the vocabulary.

After obtaining sentence-level uncertainty estimation, one can further calculate the passage-level uncertainty score by taking the average over all sentence-level uncertainty scores. In this study, we use the metrics discussed above as the **single-inference uncertainty estimation methods** and denote them as **Max Prob**, **Average Prob**, **Max Ent** and **Average Ent**.



### 3.3 Multi-inference Uncertainty Estimation

Multi-inference uncertainty estimation methods leverage the stochastic in either a model’s parameters (e.g., *Bayesian methods*) or data (e.g., *test-time data augmentation methods*) to collect a set of non-deterministic predictions. A model’s prediction uncertainty is then estimated as the divergence among those predictions.

#### 3.3.1 Metrics

Two metrics were widely used to measure such divergence: (1) variation ratio (VR) and (2) variation ratio for original prediction (VRO) [77], [91], [44]. Originally, both metrics are defined for a classification problem. Wang et al. [92] extend the definitions of VR/VRO and show that they are still effective in tasks other than classification. We introduce these two metrics in the following:

$$VR = 1 - \frac{\sum_{i=1}^T w * \frac{\sum_{j=1, j \neq i}^T (1 - \text{dist}(p_i, p_j))}{T-1}}{T}, \quad (6)$$

$$VRO = 1 - \frac{\sum_{i=1}^T (1 - \text{dist}(p_i, p_{LM}))}{T}, \quad (7)$$

where  $T$  is the number of inferences,  $\text{dist}(\cdot)$  denotes the distance function between two outputs.  $p_i$  and  $p_j$  are the inference result at  $i$ th and  $j$ th inference.  $p_{LM}$  is the prediction result from the original model  $M$ .  $w$  denotes a weight matrix.




Fig. 5: An example of multiple inferences with LLMs.

We now use an example to demonstrate the calculation of VR/VRO in Fig. 5. Suppose we have three different generations given stochasticity in either model or data.  $(1 - BLEU)$  score [93] is used as  $\text{dist}(\cdot)$  to measure the differences between two sentences<sup>1</sup>. Suppose we use identical weights for VR, then VR in this case is

$$VR = 1 - \frac{\frac{0.69+1}{2} + \frac{1+1}{2} + \frac{0.69+1}{2}}{3} \approx 0.10. \quad (8)$$

Similarly, VRO is

$$VRO = 1 - \frac{0.69 + 0.69 + 1}{3} \approx 0.21. \quad (9)$$

#### 3.3.2 Stochastic inference

To enable the stochasticity of an ML model, one popular way is to execute it with the dropout layer(s) enabled at test-time [77]. However, this is not feasible given that we focus on “black-box” analysis. We introduce two different ways to enable stochasticity in our study: (1) **sample-based method** (*Bayesian methods*) and (2) **perturbation-based method** (*test-time augmentation methods*).

**Sample-based method.** Randomness exists in LLMs’ generation process. Specifically, by controlling the parameter

*temperature* ( $t$ ), an LLM’s generation can be either deterministic or non-deterministic. When  $t$  is 0, an LLM will always choose the token with the largest probability (i.e., greedy decoding). Thus the generation will be deterministic. When  $t > 0$ , an LLM will randomly choose a token as long as its probability is larger than a threshold. A higher  $t$  will lead to a larger randomness in an LLM’s generation. In this study, we set  $t > 0$  to enable an LLM’s stochasticity and generate non-deterministic outputs. We refer to this method as *sample-based method*. We denote the usages of two different metrics as **Sample VR** and **Sample VRO**. We detail the choices of  $t$  in Sec. 4.4.

**Perturbation-based method.** An LLM’s stochasticity could also be enabled when changing the input prompt or perturbing a generated token. This also refers to the *test-time augmentation methods*. Intuitively, along the chain of the token generation, any perturbation of a generated token can affect the proceeding tokens’ generation, and possibly result in two semantically different texts. We refer to an example in Fig. 3 (C), where LLM answers the question, “How long did the Roman Empire last?” Changing the word *over* in the generation results could yield a completely different result, “lasted for *centuries*.” This phenomenon shows an intriguing property of LLMs: their stochastic nature along the prediction chain with respect to perturbation.

Therefore, in this study, we propose an adaptive method to perturb “critical” token(s) during an LLM’s generation process. Prior work has shown that tokens with high entropy in a sentence usually convey more information from the information theory perspective [94]. Consequently, perturbations at such points could possibly lead to a more pronounced change compared with others. To explore the effectiveness of using such perturbation for measuring uncertainty, we define three types of interest points to perturb: (1) the point with the highest entropy; (2) the point with the lowest entropy; (3) the point that gains the maximum entropy from the previous token in a response. In our study, we select the token that matches one of these three types of points and replace it with  $k$  other tokens with the highest probabilities. We name this method *perturbation-based method* and denote six different variants as: **Max VR**, **Max VRO**, **Min VR**, **Min VRO**, **MaxDiff VR**, and **MaxDiff VRO**. We justify the choice of  $k$  in Sec. 4.4.

## 4 STUDY DESIGN

In this section, we introduce our study design and methodology to answer our research questions in Sec. 1.

### 4.1 Subject LLMs

In order to provide a thorough evaluation of the effectiveness of uncertainty measurement as the risk assessment for LLMs in both natural language and domain-specific (i.e., code-generation) problems, we have chosen a broad spectrum of representative models. We select these models considering their availability, diversity, and computational requirements. Three open-source general LLMs, two closed-source general LLMs, and three LLMs specified for code generation are selected. We present our selected LLMs in Table 1.

1. A higher  $(1 - BLEU)$  indicates a lower similarity between two texts.

TABLE 1: Subject LLMs in our study.

LLMs	GPT-2-XL [60]	LLaMA-7B [5]	LLaMA2-7B [6]	GPT-3 [47]	GPT-3.5 [88]	InCoder [33]	SantaCoder [34]	CodeGen-6.1B [35]	Code Llama [2]
<b>Model Size</b>	1.5B	6.7B	6.7B	6.7B	175B	6.7B	1.1B	6.1B	6.7B
<b>Training Data</b>	40 GB	1,000B tokens	2T tokens	570 GB	Unknown	159 GB	268 GB	217 GB	620B tokens
<b>Domain</b>	General	General	General	General	General	Code	Code	Code	Code
<b>Provider</b>	OpenAI	MetaAI	MetaAI	OpenAI	OpenAI	MetaAI	BigCode	Salesforce	MetaAI
<b>Access</b>	open-source	open-source	open-source	closed-source	closed-source	open-source	open-source	open-source	open-source

TABLE 2: The collected NLP and programming tasks.

Dataset	Task Domain	Size	$dist(\cdot)$ metrics in VR/VRO
<b>Eli5-Category</b>	Question Answering	5,411	F1 score [95]
<b>Wiki-QA</b>	Question Answering	243	F1 score [95]
<b>CNN/Daily Mail</b>	Text Summarization	11,490	ROUGE-L [96]
<b>WMT 2014</b>	Machine Translation	3,003	BLEU [93]
<b>MBPP</b>	Code Generation	500	CodeBLEU [97]
<b>HumanEval</b>	Code Generation	164	CodeBLEU [97]

## 4.2 Tasks

To comprehensively understand the effectiveness of selected uncertainty estimation methods, we select a set of diverse and challenging tasks as a benchmark (Table. 2). Specifically, our evaluation covers four NLP tasks from three different domains (i.e., question answering, text summarization, and machine translation). Furthermore, our evaluation also includes two different code-generation tasks. Below we introduce different domains’ tasks.

**Question Answering (QA)** requires an LLM with capabilities of understanding users’ intentions, extracting learned knowledge, and organizing the responses. In this study, we select two different benchmarks for QA: *Eli5-category* [98] and *Wiki-QA* [95]. *Eli5-category* is a collection of 5,411 questions and answers gathered by and obtained from pushshift.io. *Wiki-QA* is a collection of questions extracted from *Bing* query logs with Wikipedia summaries as answers. We select 243 instances after the de-duplication and removal of questions without answers.

**Text Summarization** aims to condense a long text into a concise short summary. Different from QA, text summarization focuses on benchmarking an LLM’s capabilities of extracting critical information from a long paragraph of text. We use CNN/Daily Mail dataset [99] as the benchmark for text summarization. It comprises 11,490 news articles from CNN and the Daily Mail, with summaries obtained through the concatenation of highlighted sentences as composed by the original authors. Note that during the evaluation, we add a prompt `TL;DR` after the input text to enable an LLM’s in-context learning ability for summarization [60].

**Machine Translation** is another fundamental task to benchmark language models. In this study, we use WMT 2014 dataset [100], including 3,003 pairs of French-English transcripts, to evaluate LLMs. Similar to text summarization, We used the prompt template from the GPT-3 paper [47], where a random example translation is presented in the input for in-context learning.

**Code Generation** requires an LLM to understand both natural language (e.g., task description) and programming languages (e.g., formal structures and precise syntax). We select two datasets as benchmarks for code generation: HumanEval [101] and MBPP [102]. HumanEval consists of 164

programming problems with manually-written test cases released by OpenAI. MBPP includes 1,000 crowd-sourced Python programming problems with entry-level difficulties.

## 4.3 Evaluation Metrics

**NLP Tasks.** We use *semantic distance* to measure the performance of LLMs’ generation for NLP tasks. Specifically, we first embed the text using sentence-transformer [103] and compute the cosine distance of the embeddings. A higher cosine distance value indicates a greater level of similarity. This metric can generalize to different NLP tasks regardless of the length and form of generation. Note that while there are some other metrics specialized for each NLP task, which relies on string-matching (e.g., F1 score for text summarization), we argue two major shortcomings exist when using such metrics. First, string-matching might not accurately capture the divergence between the model’s output and the ground truth (e.g., when the output and ground truth are semantically equivalent while lexically different). Additionally, LLMs without fine-tuning might generate responses in a free-form manner that are often longer than the ground truth in, e.g., text summarization and question answering. In this case, even if an LLM answers correctly, metrics based on string-matching could still mis-indicate its performance due to the length differences.

**Code Generation Task.** Different from NLP tasks, it is relatively easy to assess the quality of generated code. We introduce a quality score  $Q$  given by  $Q = (Q_{syntax} + Q_{semantics})/2$ .  $Q_{syntax} = 1$  if the generated code is syntactically correct. Meanwhile,  $Q_{semantics}$  is the proportion of test cases the generated code passes.

## 4.4 Experiment Settings

**Uncertainty measurement.** We investigate twelve uncertainty estimation methods in our experiments: four *single-inference* methods, two *sample-based* methods, and six *perturbation-based* methods. We set the number of inferences  $T$  to 5 for both *sample-based* and *perturbation-based* methods. This is for a fair comparison between these two methods since closed-source LLMs (i.e., GPT-3, GPT-3.5) only provide access to top-5 tokens for each token’s generation. Therefore, only 5 inferences can be obtained with the *perturbation-based* method on closed-source LLMs. For temperature  $t$ , we follow the previous work and set  $t$  to 0.7 when enabling on an LLMs’ stochasticity [104], [105].

For distance metric in VR and VRO (i.e.,  $dist(\cdot)$  in Eq. 6&7), we consider both (1) general metrics based on embeddings’ cosine distance and (2) task-specific metrics. For general metrics, we use *all-mpnet-base-v2* [106] for natural language tasks and *codebert-base* [57] for code data. For

TABLE 3: Pearson correlation coefficients between uncertainty scores and LLMs’ performance on four NLP tasks. The results of VR/VRO are presented as “*cosine distance-based (task-specific distance-based)*.” Highest correlations from different categories are ranked and highlighted as *top-1*, *top-2*, and *top-3*.

Dataset	LLM	Single-inference Method				Sample-based Method	
		Max Prob	Average Prob	Max Ent	Average Ent	Sample VR	Sample VRO
CNN/Daily Mail	GPT2	-0.168	-0.198	-0.036	-0.158	-0.325(-0.091)	-0.510 (-0.227)
	LLaMA	-0.225	-0.420	-0.086	-0.415	-0.380(-0.229)	-0.520(-0.337)
	LLaMA2	-0.096	-0.109	-0.139	-0.110	-0.370(-0.260)	-0.534(-0.294)
	GPT3	-0.231	-0.202	-0.145	-0.170	-0.229(-0.100)	-0.394(-0.244)
	GPT3.5	-0.119	-0.119	-0.036	-0.106	-0.223(-0.131)	-0.220(-0.158)
Eli5-Category	GPT2	-0.021	0.021	-0.077	0.014	-0.333(-0.168)	-0.567 (-0.335)
	LLaMA	0.025	0.030	-0.040	0.007	-0.257 (-0.136)	-0.655 (-0.396)
	LLaMA2	-0.000	-0.008	0.008	-0.018	-0.196(-0.006)	-0.300(0.024)
	GPT3	-0.054	-0.207	-0.016	-0.172	-0.236(-0.084)	-0.382(-0.140)
	GPT3.5	-0.037	-0.240	-0.084	-0.289	-0.168(-0.021)	-0.223(-0.046)
Wiki-QA	GPT2	-0.110	-0.003	-0.134	-0.004	-0.433(-0.150)	-0.682 (-0.434)
	LLaMA	-0.057	0.033	-0.099	-0.020	-0.347(-0.260)	-0.598(-0.468)
	LLaMA2	-0.181	-0.193	-0.166	-0.140	-0.281(-0.206)	-0.374(-0.184)
	GPT3	-0.112	-0.192	0.017	-0.107	-0.326(-0.061)	-0.376(-0.139)
	GPT3.5	-0.051	-0.107	-0.065	-0.105	-0.401(-0.085)	-0.425(-0.079)
WMT 2014	GPT2	0.239	0.073	0.282	0.081	-0.556 (-0.253)	-0.656 (-0.276)
	LLaMA	-0.044	-0.153	-0.060	-0.166	-0.472 (-0.353)	-0.837 (-0.636)
	LLaMA2	0.157	-0.000	0.185	-0.015	0.068(0.009)	0.125(0.163)
	GPT3	-0.046	-0.162	-0.039	-0.176	-0.250(-0.254)	-0.251(-0.339)
	GPT3.5	0.089	-0.158	0.050	-0.175	-0.244(-0.065)	-0.233(-0.017)

Dataset	LLM	Perturbation-based					
		Max VR	Max VRO	Min VR	Min VRO	MaxDiff VR	MaxDiff VRO
CNN/Daily Mail	GPT2	0.407 (0.448)	0.356(0.430)	0.260(0.269)	0.376(0.403)	0.400(0.436)	0.397(0.439)
	LLaMA	0.447 (0.504)	0.410(0.484)	0.324(0.335)	0.443(0.495)	0.420(0.469)	0.429(0.479)
	LLaMA2	-0.204(-0.161)	-0.388(-0.154)	-0.075(-0.055)	-0.418(-0.226)	-0.194(-0.152)	-0.378(-0.143)
	GPT3	-0.096(-0.084)	-0.233(-0.209)	-0.001(-0.049)	0.034(0.041)	-0.119(-0.121)	-0.228(-0.182)
	GPT3.5	0.064(-0.086)	-0.024(-0.087)	-0.116(-0.176)	-0.068(-0.177)	-0.042(-0.138)	0.072(0.021)
Eli5-Category	GPT2	-0.251(-0.158)	-0.414(-0.247)	-0.055(-0.081)	-0.121(-0.186)	-0.268(-0.174)	-0.473(-0.284)
	LLaMA	-0.163(-0.100)	-0.442(-0.258)	-0.143(-0.140)	-0.102(-0.142)	-0.172(-0.073)	-0.357(-0.178)
	LLaMA2	-0.079(0.008)	-0.203(-0.017)	-0.086(0.002)	-0.218(0.026)	-0.086(0.018)	-0.210(-0.011)
	GPT3	0.168(0.179)	0.026(0.079)	-0.236(-0.246)	-0.066(-0.135)	0.140(-0.003)	0.184(0.047)
	GPT3.5	0.003(-0.076)	0.065(0.049)	-0.087(-0.157)	0.072(0.056)	0.160(0.092)	0.183(0.125)
Wiki-QA	GPT2	-0.441(-0.284)	-0.617(-0.373)	-0.169(-0.186)	-0.136(-0.145)	-0.270(-0.107)	-0.557(-0.279)
	LLaMA	-0.200(-0.126)	-0.478(-0.416)	-0.090(-0.206)	-0.135(-0.348)	-0.227(-0.257)	-0.338(-0.404)
	LLaMA2	-0.057(0.050)	-0.081(0.014)	-0.064(-0.001)	-0.164(-0.036)	-0.077(0.023)	-0.111(-0.006)
	GPT3	-0.064(0.049)	-0.204(-0.186)	-0.042(0.048)	-0.247(-0.146)	-0.042(-0.020)	-0.095(-0.111)
	GPT3.5	0.092(0.065)	0.078(0.056)	-0.017(-0.113)	-0.030(-0.109)	0.157(0.014)	0.214(0.065)
WMT 2014	GPT2	-0.025(0.067)	0.042(0.071)	-0.059(0.009)	0.020(0.077)	-0.093(-0.010)	-0.041(0.007)
	LLaMA	0.153(0.086)	0.099(-0.153)	0.094(0.055)	0.155(0.009)	-0.038(-0.033)	-0.012(-0.150)
	LLaMA2	-0.022(-0.027)	-0.036(0.058)	-0.064(-0.029)	-0.190(0.024)	-0.026(-0.028)	-0.040(0.058)
	GPT3	-0.076(0.006)	0.016(0.021)	-0.033(0.095)	0.215(0.138)	-0.103(0.001)	0.013(0.040)
	GPT3.5	-0.277(-0.363)	-0.101(-0.081)	-0.099(-0.048)	-0.007(-0.003)	-0.105(-0.201)	-0.031(-0.024)

task-specific metrics, we consider the following choices in Table 2.

**Experiment setups.** Due to the API constraints, an inference of closed-source LLMs (i.e., GPT3 and GPT3.5) could take up to 20 seconds. Consequently, evaluating an LLM on CN-N/Daily Mail dataset with 11,490 instances would require nearly 575 hours. Therefore, we randomly sample each NLP task’s dataset with 100 instances when evaluating closed-source LLMs. We further sample 40 out of 164 and 125 out of 500 instances for HumanEval and MBPP, respectively. For closed-source models, we evaluated all instances.

**Hardware and software dependencies.** To conduct our large-scale experiments, we utilize a cluster of servers with AMD 3955WX CPU (3.9GHz), 256GB RAM, and four NVIDIA A4000 GPUs (16GB VRAM of each). The evaluation of open-source LLMs takes more than 864 GPU hours.

## 5 RESULTS

### 5.1 RQ1: Uncertainty Estimation for NLP Tasks

To answer this research question, we evaluate five general LLMs. We present the results of Pearson correlation coefficients between uncertainty scores and LLMs’ performance in Table 3. Such correlation is an indicator of whether the uncertainty estimation can predict LLMs’ performance and further perform the risk assessment. The higher the absolute value of the coefficient, the stronger the correlation. We investigate the results on NLP tasks from three perspectives:

**Uncertainty Measurement Techniques.** As can be observed from Table 3, uncertainties estimated via *sample-based* methods generally yield the highest correlation to an LLM’s performance. Specifically, sample-based VRO (**Sample VRO**) achieves the best performance in sixteen out of twenty cases (5 LLMs  $\times$  4 tasks). In contrast, *single-inference* methods exhibit relatively low performance. Only 5 out of 80 scenarios have an absolute correlation value greater than 0.25. This indicates that relying on single-inference uncertainty



estimation might be unreliable in practical applications without a more refined strategy. An extreme case can be when using LLaMA for the machine translation task, where sample-based VRO has significantly better performance (-0.837 vs. -0.166) compared with the best single-inference method (i.e., average entropy). However, even for the most promising sample-based methods, the majority correlation is below -0.6 and their performance can still further space for enhancement (e.g., -1.00 correlation), calling for the design of more advanced techniques.

**Finding 1:** *Sample-based VRO* achieves the best performance for five different LLMs on NLP tasks in most cases, surpassing single-inference methods by a large margin. While its potential is evident, further enhancement with advanced technique design is needed and promising for industrial deployment.

*Perturbation-based* methods demonstrate moderate effectiveness, exceeding single-inference metrics in a substantial proportion of tasks. For the selection of perturbation points, the maximum entropy point is usually better than the other two. Surprisingly, the minimum entropy point can sometimes work well, especially for GPT-3 and Llama2. For GPT-3, it might be because we can not obtain the probability lists across the entire vocabulary, and the estimated entropy is biased. For Llama2, we observed that it is substantially more “certain” than Llama-1. The 75th percentile of the entropy of all tokens generated by Llama2 is nearly 0, with a mean value of 0.18. Conversely, Llama-1 has a value of 2.62 at the 75th percentile, with a mean of 1.52. As such, the entropy selection strategy might not work well on Llama2.

Another interesting observation is that when using GPT-2 and LLaMA on the CNN/Daily Mail dataset (text summarization), the highest Pearson correlation coefficients of perturbation-based methods are positive. The results indicate that, a higher uncertainty may even yield a better performance in such cases. Upon several case studies, we find that this observation could be attributed to the long input in the text summarization dataset. Both GPT-2 and LLaMA may be out of focus on the input context. Therefore, the LLMs do not understand the summarization task and give similar responses despite the stochasticity from perturbations, resulting in a low uncertainty score. On the contrary, if an LLM understands the task description at the end of the prompt, it will respond to the perturbation more severely, causing a higher uncertainty.

**Finding 2:** Perturbation-based methods are more inclined to produce model-specific outcomes. Stakeholders might need to perform model-specific optimization if possible.

**Influence of Distance Functions.** We also find that task-specific distance does not provide better results compared with cosine distance between embeddings in both *sample-based* and *perturbation-based* methods. On average, using cosine distance achieves an increase of 0.093 for *perturbation-based* methods and an increase of 0.189 for *sample-based* methods across different tasks and LLMs.




Fig. 6: Performance of LLMs on NLP tasks

**Finding 3:** The cosine distance function yields a better performance on NLP tasks when leveraging stochasticity to estimate an LLM’s prediction uncertainty.

**Influence of Models.** Combining LLMs’ performance in Figure 6 and uncertainty methods’ effectiveness in Table 3, we can observe that it is challenging for these uncertainty methods to assess the risks for models that are either *exceptionally good* (e.g., GPT-3.5) or *notably poor* (e.g., Llama2). This might be because the current evaluated methods struggle to discern the nuanced differences attributed to aleatoric uncertainty [27] (knowledge deficiency). Most of the results are likely dominated by epistemic uncertainty [27] (inherent randomness), leading to a decline in detection performance. In other words, when generated contents are unanimously good or bad, the computed uncertainty is more likely driven by inherent randomness rather than variations in the input data.

**Finding 4:** Uncertainty-based risk assessments for exceptionally good or notably poor LLMs appear to be limited. The former is especially prevalent in closed-source commercial models used in real-world applications. When an LLM’s performance is very high, analyzing its potential issues can be intrinsically difficult, calling for more advanced methods that can work even with limited information.

## 5.2 RQ2: Limitations in Uncertainty Estimation for NLP Tasks

In this RQ, we further explore two limitations of evaluated uncertainty estimation methods.

**Influence of prompt.** An intriguing anomaly in RQ1 is Llama2, which underperforms significantly on summarization and translation, and its uncertainty estimation for all four tasks is ineffective. This is counter-intuitive. After manually investigating some randomly chosen samples, we found this might be because of the input prompt. Using a fine-tuned chat version of Llama2, we incorporated a system prompt template shown in the original paper, starting with

You are a helpful, respectful, and honest assistant ...

However, this prompt template has three downsides. We will use the WMT-14 translation task as examples: (1) The LLM will sometimes directly reject to respond, offering something such as “ I cannot provide a translation for

TABLE 4: AUC scores for detecting erroneous code generated by LLMs. The results of VR/VRO are presented as “*cosine distance-based (task-specific distance-based)*.” Highest scores are ranked and highlighted as *top-1*, *top-2*, and *top-3*.

Dataset	LLM	Single-inference Method				Sample-based Method	
		Max Prob	Average Prob	Max Ent	Average Ent	Sample VR	Sample VRO
HumanEval	Codegen	0.728	0.625	0.694	0.604	0.746(0.786)	0.821(0.825)
	Incoder	0.634	0.632	0.633	0.629	0.758(0.760)	0.799(0.778)
	Santacoder	0.830	0.685	0.630	0.756	0.556(0.639)	0.843(0.710)
	GPT3.5	0.500	0.590	0.383	0.490	0.630(0.667)	0.607(0.587)
	Code Llama	0.744	0.618	0.752	0.601	0.697(0.681)	0.731(0.749)
MBPP	Codegen	0.482	0.396	0.470	0.406	0.533(0.689)	0.586(0.705)
	Incoder	0.409	0.488	0.420	0.493	0.569(0.699)	0.665(0.695)
	Santacoder	0.671	0.565	0.659	0.442	0.558(0.692)	0.631(0.817)
	GPT3.5	0.582	0.556	0.606	0.560	0.592(0.563)	0.561(0.576)
	Code Llama	0.499	0.418	0.483	0.433	0.578(0.668)	0.603(0.705)

Dataset	LLM	Perturbation-based					
		Max VR	Max VRO	Min VR	Min VRO	MaxDiff VR	MaxDiff VRO
HumanEval	Codegen	0.484(0.540)	0.480(0.522)	0.461(0.470)	0.655(0.604)	0.525(0.548)	0.560(0.607)
	Incoder	0.579(0.704)	0.606(0.689)	0.342(0.424)	0.586(0.450)	0.566(0.578)	0.616(0.647)
	Santacoder	0.574(0.315)	0.716(0.735)	0.682(0.565)	0.586(0.747)	0.590(0.327)	0.762(0.765)
	GPT3.5	0.500(0.503)	0.457(0.560)	0.497(0.503)	0.420(0.607)	0.533(0.533)	0.393(0.477)
	Code Llama	0.597(0.613)	0.668(0.644)	0.498(0.555)	0.643(0.582)	0.609(0.538)	0.725(0.680)
MBPP	Codegen	0.592(0.619)	0.560(0.627)	0.499(0.515)	0.522(0.547)	0.613(0.638)	0.543(0.616)
	Incoder	0.611(0.686)	0.650(0.688)	0.434(0.565)	0.593(0.542)	0.694(0.704)	0.678(0.695)
	Santacoder	0.578(0.549)	0.429(0.659)	0.496(0.473)	0.278(0.493)	0.597(0.574)	0.476(0.655)
	GPT3.5	0.587(0.628)	0.511(0.544)	0.550(0.495)	0.508(0.511)	0.498(0.571)	0.486(0.528)
	Code Llama	0.586(0.630)	0.521(0.621)	0.454(0.507)	0.501(0.568)	0.622(0.679)	0.566(0.667)

that statement as it is not factually coherent.” (2) The LLM misinterprets examples in the in-context prompt, translating the (fixed) example itself. This leads to duplicate translations and results in poor performance in translation. (3) Many of the responses will begin with greetings such as “Thank you for your kind and respectful instructions!” Each of these three scenarios can negatively affect the calculation of multi-inference uncertainty, as the fine-tuned responses from the LLM, to some extent, surpass inherent randomness.

To study the prompt influence further, we removed the system prompt and re-run all the experiments on Llama2. The results are shown in Table 5. Overall, there’s a significant increase in all correlations, encompassing the original two high-performing QA tasks. When comparing the top-1 scores across the four tasks, the average increase in the absolute value of correlation coefficients is 0.272. The most notable improvement is observed in the WMT-14 dataset, where the correlation shifts from -0.190 to -0.746. Given that RLHF-based prompt tuning is pivotal for SOTA chat models, such as open-sourced Llama2 and commercial ChatGPT, it might be better for future work to take this into consideration when designing new methods.

**Finding 5:** Prompts could significantly influence uncertainty estimation in some cases. Prompt templates can lead LLMs to exhibit different behaviors, compromising the accuracy of uncertainty estimation.

**Relation of uncertainty and inaccuracy.** We find that a low uncertainty does not guarantee that an LLM’s response is reliable. Specifically, an LLM can generate highly confident responses while with non-factual information.

A representative example can be observed when GPT-2 attempts to respond to a question from the Eli5 category dataset:

Why do scientists research commonly known stuff?

The LLM misunderstands this question and only captures the words “scientists” and “commonly”. Consequently, it turns towards a discussion about the human body parts that scientists commonly possess, delivering responses such as “We all have the same DNA, we have the same genes, we all have the same ears...” In this case, an LLM consistently generates similar responses related to body parts despite introducing stochasticity through perturbation. Even though the responses are completely incorrect, an LLM still yields a low uncertainty score.

Conversely, a higher degree of uncertainty also does not necessarily imply that an LLM’s prediction is incorrect. We show an example that GPT 3.5 summarizes a piece of news from the CNN/Daily Mail dataset. In this case, the ground truth summary includes two details: *a fire occurrence at a park* and *the absence of injuries*. For sample-based uncertainty measurement, all five generated samples incorporate these two pieces of information but also furnish additional varied information, such as the park’s owner and its intended use. Such extra information further leads to a higher variance in the generated samples’ embedding, resulting in a high uncertainty score despite that the LLM’s prediction is reliable.

**Finding 6:** Uncertainty is not always correlated to inaccuracy. Future work may also consider combining other features or indicators (besides uncertainty) to the risk assessment for better performance.

### 5.3 RQ3: Uncertainty Estimation for Code Generation

To answer this research question, we evaluate five LLMs designed for code tasks. Their performances are shown

TABLE 5: Pearson correlation coefficients of LLaMA2 without system prompts on four NLP tasks. The results of VR/VRO are presented as “*cosine distance-based (task-specific distance-based)*.” Highest correlations from different categories are ranked and highlighted as *top-1*, *top-2*, and *top-3*.

Dataset	Single-inference Method				Sample-based Method	
	Max Prob	Average Prob	Max Ent	Average Ent	Sample VR	Sample VRO
CNN/Daily Mail	-0.349	-0.480	-0.435	-0.528	<b>-0.601</b> (-0.528)	<b>-0.638</b> (-0.522)
Eli5-Category	-0.054	-0.316	-0.182	-0.362	-0.354(-0.020)	<b>-0.444</b> (0.006)
Wiki-QA	-0.093	-0.122	-0.073	-0.173	<b>-0.534</b> (-0.212)	<b>-0.589</b> (-0.138)
WMT 2014	-0.006	0.176	-0.145	-0.003	<b>-0.611</b> (-0.527)	<b>-0.746</b> (-0.578)

Dataset	Perturbation-based					
	Max VR	Max VRO	Min VR	Min VRO	MaxDiff VR	MaxDiff VRO
CNN/Daily Mail	0.098(-0.054)	-0.006(-0.053)	0.000(-0.087)	-0.332(-0.320)	0.123(-0.037)	0.025(-0.028)
Eli5-Category	-0.191(-0.124)	<b>-0.513</b> (-0.293)	-0.183(-0.042)	-0.256(0.059)	-0.169(-0.109)	<b>-0.498</b> (-0.276)
Wiki-QA	0.007(0.029)	-0.008(0.070)	-0.193(-0.123)	<b>-0.438</b> (-0.103)	0.049(0.010)	-0.058(0.011)
WMT 2014	0.034(-0.066)	0.230(-0.032)	-0.400(-0.409)	<b>-0.608</b> (-0.515)	0.030(-0.055)	0.266(0.014)

in Figure 7. When evaluating the efficacy of uncertainty methods in identifying erroneous code, we treat the problem as a binary detection task. Code is labeled as 0 if completely correct (e.g.,  $Q = 1$ ) and 1 if not. We take the uncertainty scores as indicators, and higher uncertainty suggests a greater likelihood of errors in the code. Subsequently, we calculate Area Under the Receiver Operating Characteristic Curve (AUC) scores for each method. The results are shown in Table 4.

**Uncertainty Measurement Techniques.** Even with the distinct task of code generation and the different evaluation metric of AUC, we still observe a similar trend on code generation tasks than on NLP tasks, i.e., *sample-based* methods dominate across two different datasets and different LLMs. A nuanced distinction is that perturbation-based methods are less effective in pinpointing erroneous code. In NLP tasks, these methods account for 38.3% (23/60) of the top-3 rankings across all settings. In contrast, they represent only 26.7% (8/30) of the top-3 rankings in code generation tasks.

**Finding 7:** Sample-based methods are still the most effective in risk assessment for code generation.

**Influence of Distance Functions.** Different from the results on NLP tasks, the cosine distance between code embeddings under-perform significantly in comparison with the task-specific distance function, CodeBLEU. Task-specific metric prevails in 20 out of 23 cases for the TOP-3 results.

**Finding 8:** Distance function can play an important role in uncertainty estimation as indicated by the result difference between domains of NLP and code. A more carefully designed distance function that suits downstream tasks could potentially enhance the risk assessment effectiveness.

**Influence of Models.** Among all models, GPT-3.5 is the most challenging model to assess code quality using uncertainty estimation. The best AUC score is 0.667 on the HumanEval dataset and 0.628 on the MBPP dataset. Both of them are below 0.7. According to the violin plots in Figure 7, GPT-3.5 performs best, and all the code it generates is syntactically correct, with a significant amount of code passing at least




Fig. 7:  $Q$  scores of LLMs on code generation tasks

one test case. Identifying errors in such “partially correct” code poses a greater challenge. This aligns with our previous finding in NLP tasks.

#### 5.4 RQ4: Limitations in Uncertainty Estimation for Code Generation

**Limitations of detecting subtle errors.** In RQ3, we hypothesize that the relatively poor performance of uncertainty estimation in detection GPT-3.5’s errors is due to the challenges in pinpointing subtle mistakes that existed in the generated code. To validate this hypothesis, we conduct additional experiments in this RQ. We evaluate all the methods across four LLMs on two datasets under two settings. In setting *A*, only syntactically correct code is kept (e.g.,  $Q \geq 0.5$ ). In setting *B*, only syntactically incorrect and completely correct code is kept (e.g.,  $Q = 1$  or  $Q < 0.5$ ). We exclude GPT-3.5 since all its code is syntactically correct. We record the best AUC score for each model under both settings. We leave the full result on our website.

Comparing the performances between setting *A* and setting *B*, we observe that, with the exception of SantaCoder—which experiences a drop from 0.878 in setting *A* to 0.843 in setting *B* on HumanEval—all other methods show a notable improvement. The average score increase is 0.12, with CodeGen achieving the highest jump of 0.183. SantaCoder’s decline might be attributed to its frequent generation of syntactically incorrect code, making the partially correct code an easily detectable outlier.

TABLE 6: The average distances between generated code.

	Perturbation		Sample	
	Open-source	Closed-source	Open-source	Closed-source
Cosine	0.03±0.04	0.03±0.05	0.03±0.02	0.02±0.01
CodeBLEU	0.53±0.19	0.45±0.19	0.71±0.10	0.41±0.16

**Finding 9:** Current uncertainty methods are more suitable for detecting obvious errors instead of subtle errors.

**Limitations of distance functions.** Different from our observations in RQ1, cosine distance metric does not show dominant performance compared with task-specific metric on estimating LLMs’ risks when generating code. This indicates that it is non-trivial for selected embedding model [57] to detect minor code differences due to an LLM’s stochasticity.

The efficiency of multi-inference metrics heavily depends on the precise estimation of the distance between data points. Both RQ1 and RQ3 highlight the crucial role of the distance function. However, it appears that the cosine distance metric falls short in code-related tasks. It is non-trivial for current embedding LLMs to detect subtle differences generated by randomness or perturbations.

We calculate pairwise distances within each group of responses generated via sample-based and perturbation-based methods. As shown in Table 6, the average cosine distance is significantly lower than its task-specific counterpart (i.e., CodeBLEU [97]). This directly affects the effectiveness of VR/VRO when determining the degree of uncertainty.

Nevertheless, CodeBLEU [97] also comes with its limitations. For instance, two programs that only differ from variable names might lead to a large CodeBLEU distance, resulting in erroneous uncertainty estimation.

**Finding 10:** Neither cosine distance nor CodeBLEU could accurately assess the difference between two programs, resulting in under-performed uncertainty estimation for LLMs compared with NLP tasks. Characterizing the true difference between code can be challenging and is a limitation to performing accurate uncertainty estimation.

## 6 IMPLICATION AND OPPORTUNITY

**Ask more, get more.** In our study, *multi-inference* methods perform better than single-inference methods across different tasks in the most cases. As for a *black-box* LLM, getting a comprehensive understanding beforehand or only through a single deterministic inference could be challenging. Instead, the more we query an LLM, the clearer we can get about its internal knowledge regarding a specific aspect. We hypothesize that this is because by querying models multiple times, we gain more knowledge from them. This might be a promising technique when the model is black-box.

**Prompt is important.** Prompt has long been proven to be a key factor in LLM’s performance. In our study, we further demonstrate its potential to significantly affect uncertainty estimation’s efficacy (Finding 5). Specifically, the uncertain

behavior of LLMs might be profoundly impacted by the prompt used in RLHF [30]. Original uncertain answers may be supplanted by human-favored responses when integrating specific prompt templates. A prospective research direction is to explore the influence of the RLHF process on uncertainty estimation and to discern strategies for more accurate estimations, both from training (e.g., calibrate the model better) and inference (e.g., refine estimation methods) standpoints. Another interesting perspective is to design a better prompt (e.g., instruct LLMs to switch to the uncertainty estimation mode) to enhance the precision of the measurement.

### Model-specific uncertainty estimation might be beneficial.

Although all the methods chosen in the paper are black-box, we observed considerable variations in their effectiveness across different models (e.g., Finding 2). Thus, to enhance the efficacy of an uncertainty-based risk assessment system, stakeholders might need to tailor their methods and undertake model-specific optimizations. These adjustments could be necessary even between different model versions (e.g., Llama and Llama2).

**Subtle errors can be hard to detect.** We observed that selected methods could struggle to detect subtle errors in partially correct code and are easier to obvious errors (Finding 9). One possible future direction is to improve their sensitivity by separating the uncertainty caused by model inability and that stemming from inherent randomness. Another direction could involve constructing a multi-stage system, with uncertainty-related methods at the forefront, followed by other techniques (e.g., white-box).

**Better perturbation strategy is needed for more accurate uncertainty estimation.** Our *perturbation-based* methods leverage the unique characteristic of an autoregressive language model to perturb its decoding process, which shows moderate uncertainty estimation performance in the experiments. Compared with *sample-based* methods, the *perturbation-based* methods do not require access and tuning the temperature setting ( $T$ ). Despite the fact that the *perturbation-based* methods underperform the *sample-based* methods in general, we believe the *perturbation-based* methods could be further improved with, e.g., a more fine-grained strategy to identify key point for perturbation.

## 7 THREATS TO VALIDITY

In terms of **internal threats**, the selection of uncertainty estimation techniques can be a threat that affects the validity of our findings. In our study, we tried our best and collect as many as 12 existing uncertainty estimation methods from different categories (single deterministic, Bayesian, and test-time augmentation), to better understand the effectiveness of uncertainty estimation under the scope of LLMs’ erroneous generations.

In terms of **external threats**, the subject LLM selection for the evaluation could be another threat. The generality of our uncertainty estimation methods to LLMs beyond the subject ones another threat since LLMs may have various uncertainty characteristics on different tasks. To mitigate this, we selected SOTA LLMs from various sources, encompassing both open-source and closed-source. These diverse LLMs

also span a wide range of tasks, from general-purpose to those fine-tuned for code-specific tasks.

## 8 CONCLUSION

This paper initiates an early exploratory study toward understanding the risk assessment of LLMs from the lens of uncertainty estimation. A large-scale evaluation of twelve different uncertainty estimation techniques on nine LLMs and six tasks is conducted. A further in-depth analysis was made to investigate the correlations between LLMs' prediction uncertainty and their performance. Understanding the potential risks of LLMs could be of great importance for industrial-scale applications. Our results confirm that uncertainty estimation can be a promising direction for potential risk assessment of LLMs in both NLP and code-generation domain tasks. However, there can still be much space and opportunity to design more advanced uncertainty estimation techniques to characterize the risks of an LLM more effectively. Moreover, other possibly useful quality indicators besides uncertainty could also be designed to better characterize the capability boundary of an LLM from multiple perspectives. With the recently increasing demand and urgency for trustworthiness assurance of LLMs in industry, we hope this paper could potentially inspire researchers and practitioners, to join the force to design novel techniques and toolchain support and together conquer many new relevant challenges. We also make the replication package of this paper available, to enable further research towards realizing trustworthy LLMs for industrial usage.

## REFERENCES

- [1] T. Markov, C. Zhang, S. Agarwal, F. E. Nekoul, T. Lee, S. Adler, A. Jiang, and L. Weng, "A holistic approach to undesired content detection in the real world," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 37, no. 12, 2023, pp. 15 009–15 018.
- [2] B. Rozière, J. Gehring, F. Gloeckle, S. Sootla, I. Gat, X. E. Tan, Y. Adi, J. Liu, T. Remez, J. Rapin *et al.*, "Code llama: Open foundation models for code," *arXiv preprint arXiv:2308.12950*, 2023.
- [3] "Chatgpt," <http://chat.openai.com>, 2023.
- [4] L. Wu, Z. Zheng, Z. Qiu, H. Wang, H. Gu, T. Shen, C. Qin, C. Zhu, H. Zhu, Q. Liu *et al.*, "A survey on large language models for recommendation," *arXiv preprint arXiv:2305.19860*, 2023.
- [5] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar *et al.*, "Llama: Open and efficient foundation language models," *arXiv preprint arXiv:2302.13971*, 2023.
- [6] H. Touvron, L. Martin, K. Stone, P. Albert, A. Almahairi, Y. Babaei, N. Bashlykov, S. Batra, P. Bhargava, S. Bhosale *et al.*, "Llama 2: Open foundation and fine-tuned chat models," *arXiv preprint arXiv:2307.09288*, 2023.
- [7] Meta, "The llama ecosystem: Past, present, and future," 2023. [Online]. Available: <https://perma.cc/YJ4E-FB95>
- [8] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Bang, A. Madotto, and P. Fung, "Survey of hallucination in natural language generation," *ACM Computing Surveys*, vol. 55, no. 12, pp. 1–38, 2023.
- [9] A. Tamkin, M. Brundage, J. Clark, and D. Ganguli, "Understanding the capabilities, limitations, and societal impact of large language models," *arXiv preprint arXiv:2102.02503*, 2021.
- [10] A. Abid, M. Farooqi, and J. Zou, "Persistent anti-muslim bias in large language models," in *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, 2021, pp. 298–306.
- [11] Y. Bang, S. Cahyawijaya, N. Lee, W. Dai, D. Su, B. Wilie, H. Love-nia, Z. Ji, T. Yu, W. Chung *et al.*, "A multitask, multilingual, multimodal evaluation of chatgpt on reasoning, hallucination, and interactivity," *arXiv preprint arXiv:2302.04023*, 2023.
- [12] D. Johnson, R. Goodman, J. Patrinely, C. Stone, E. Zimmerman, R. Donald, S. Chang, S. Berkowitz, A. Finn, E. Jahangir *et al.*, "Assessing the accuracy and reliability of ai-generated medical responses: an evaluation of the chat-gpt model," 2023.
- [13] S. Bubeck, V. Chandrasekaran, R. Eldan, J. Gehrke, E. Horvitz, E. Kamar, P. Lee, Y. T. Lee, Y. Li, S. Lundberg *et al.*, "Sparks of artificial general intelligence: Early experiments with gpt-4," *arXiv preprint arXiv:2303.12712*, 2023.
- [14] J. Schuett, N. Dreksler, M. Anderljung, D. McCaffary, L. Heim, E. Bluemke, and B. Garfinkel, "Towards best practices in agi safety and governance: A survey of expert opinion," *arXiv preprint arXiv:2305.07153*, 2023.
- [15] M. Anderljung, J. Barnhart, J. Leung, A. Korinek, C. O'Keefe, J. Whittlestone, S. Avin, M. Brundage, J. Bullock, D. Cass-Beggs *et al.*, "Frontier ai regulation: Managing emerging risks to public safety," *arXiv preprint arXiv:2307.03718*, 2023.
- [16] Microsoft, "Responsible ai standard, v2," 2022. [Online]. Available: <https://perma.cc/S466-T3HD>
- [17] OpenAI, "Our approach to ai safety," 2023. [Online]. Available: <https://perma.cc/4M6Z-LGX6>
- [18] Amazon, "Transform responsible ai from theory into practice," 2023. [Online]. Available: <https://perma.cc/V6LF-ZPD7>
- [19] Google, "Responsible ai practices," 2023. [Online]. Available: <https://perma.cc/SAG5-TEHL>
- [20] L. Koessler and J. Schuett, "Risk assessment at agi companies: A review of popular risk assessment techniques from other safety-critical industries," *arXiv preprint arXiv:2307.08823*, 2023.
- [21] Meta, "Facebook's five pillars of responsible ai," 2021. [Online]. Available: <https://ai.meta.com/blog/facebook-five-pillars-of-responsible-ai/>
- [22] —, "Responsible ai," 2023. [Online]. Available: <https://perma.cc/9BAY-NM9D>
- [23] —, "Building generative ai features responsibly," 2023. [Online]. Available: <https://about.fb.com/news/2023/09/building-generative-ai-features-responsibly/>
- [24] N. Kokhlikyan, V. Miglani, M. Martin, E. Wang, B. Alsallakh, J. Reynolds, A. Melnikov, N. Klushkina, C. Araya, S. Yan *et al.*, "Captum: A unified and generic model interpretability library for pytorch," *arXiv preprint arXiv:2009.07896*, 2020.
- [25] D. Kiela, H. Firooz, A. Mohan, V. Goswami, A. Singh, C. A. Fitzpatrick, P. Bull, G. Lipstein, T. Nelli, R. Zhu *et al.*, "The hateful memes challenge: Competition report," in *NeurIPS 2020 Competition and Demonstration Track*. PMLR, 2021, pp. 344–360.
- [26] U. Bhatt, J. Antorán, Y. Zhang, Q. V. Liao, P. Sattigeri, R. Fogliato, G. Melançon, R. Krishnan, J. Stanley, O. Tickoo *et al.*, "Uncertainty as a form of transparency: Measuring, communicating, and using uncertainty," in *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, 2021, pp. 401–413.
- [27] E. Hüllermeier and W. Waegeman, "Aleatoric and epistemic uncertainty in machine learning: An introduction to concepts and methods," *Machine Learning*, vol. 110, pp. 457–506, 2021.
- [28] O. Rahmati, B. Choubin, A. Fathabadi, F. Coulon, E. Soltani, H. Shahabi, E. Mollaefar, J. Tiefenbacher, S. Cipullo, B. B. Ahmad *et al.*, "Predicting uncertainty of machine learning models for modelling nitrate pollution of groundwater using quantile regression and uneec methods," *Science of the Total Environment*, vol. 688, pp. 855–866, 2019.
- [29] P. Manakul, A. Liusie, and M. J. Gales, "Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models," *arXiv preprint arXiv:2303.08896*, 2023.
- [30] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray *et al.*, "Training language models to follow instructions with human feedback," *Advances in Neural Information Processing Systems*, vol. 35, pp. 27 730–27 744, 2022.
- [31] X. Hou, Y. Zhao, Y. Liu, Z. Yang, K. Wang, L. Li, X. Luo, D. Lo, J. Grundy, and H. Wang, "Large language models for software engineering: A systematic literature review," 2023.
- [32] A. Fan, B. Gokkaya, M. Harman, M. Lyubarskiy, S. Sengupta, S. Yoo, and J. M. Zhang, "Large language models for software engineering: Survey and open problems," *arXiv preprint arXiv:2310.03533*, 2023.
- [33] D. Fried, A. Aghajanyan, J. Lin, S. Wang, E. Wallace, F. Shi, R. Zhong, S. Yih, L. Zettlemoyer, and M. Lewis, "InCoder: A generative model for code infilling and synthesis," in *The Eleventh International Conference on Learning Representations*, 2023. [Online]. Available: <https://openreview.net/forum?id=hQwb-lbM6EL>



- [34] L. B. Allal, R. Li, D. Kocetkov, C. Mou, C. Akiki, C. M. Ferrandis, N. Muennighoff, M. Mishra, A. Gu, M. Dey *et al.*, “Santacoder: don’t reach for the stars!” *arXiv preprint arXiv:2301.03988*, 2023.
- [35] E. Nijkamp, B. Pang, H. Hayashi, L. Tu, H. Wang, Y. Zhou, S. Savarese, and C. Xiong, “Codegen: An open large language model for code with multi-turn program synthesis,” *arXiv preprint arXiv:2203.13474*, 2022.
- [36] Q. Gu, “Llm-based code generation method for golang compiler testing,” in *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, 2023, pp. 2201–2203.
- [37] S. Kang, J. Yoon, and S. Yoo, “Large language models are few-shot testers: Exploring llm-based general bug reproduction,” in *2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE)*. IEEE, 2023, pp. 2312–2323.
- [38] C. Lemieux, J. P. Inala, S. K. Lahiri, and S. Sen, “Codamosa: Escaping coverage plateaus in test generation with pre-trained large language models,” in *International conference on software engineering (ICSE)*, 2023.
- [39] Z. Liu, C. Chen, J. Wang, X. Che, Y. Huang, J. Hu, and Q. Wang, “Fill in the blank: Context-aware automated text input generation for mobile gui testing,” in *2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE)*. IEEE, 2023, pp. 1355–1367.
- [40] R. Tian, Y. Ye, Y. Qin, X. Cong, Y. Lin, Z. Liu, and M. Sun, “Debugbench: Evaluating debugging capability of large language models,” 2024.
- [41] E. First, M. Rabe, T. Ringer, and Y. Brun, “Baldur: Whole-proof generation and repair with large language models,” in *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, 2023, pp. 1229–1241.
- [42] Y. Wei, C. S. Xia, and L. Zhang, “Copiloting the copilots: Fusing large language models with completion engines for automated program repair,” in *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, 2023, pp. 172–184.
- [43] M. Geng, S. Wang, D. Dong, H. Wang, G. Li, Z. Jin, X. Mao, and X. Liao, “Large language models are few-shot summarizers: Multi-intent comment generation via in-context learning,” 2024.
- [44] X. Zhang, X. Xie, L. Ma, X. Du, Q. Hu, Y. Liu, J. Zhao, and M. Sun, “Towards characterizing adversarial defects of deep learning software from the lens of uncertainty,” in *Proceedings of the ACM/IEEE 42nd International Conference on Software Engineering*, 2020, pp. 739–751.
- [45] J. Kim, R. Feldt, and S. Yoo, “Guiding deep learning system testing using surprise adequacy,” in *2019 IEEE/ACM 41st International Conference on Software Engineering (ICSE)*. IEEE, 2019, pp. 1039–1049.
- [46] M. Weiss and P. Tonella, “Uncertainty quantification for deep neural networks: An empirical comparison and usage guidelines,” *Software Testing, Verification and Reliability*, p. e1840, 2023.
- [47] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell *et al.*, “Language models are few-shot learners,” *Advances in neural information processing systems*, vol. 33, pp. 1877–1901, 2020.
- [48] T. Kuhn, H. Niemann, and E. G. Schukat-Talamazzini, “Ergodic hidden markov models and polygrams for language modeling,” in *Proceedings of ICASSP’94. IEEE International Conference on Acoustics, Speech and Signal Processing*, vol. 1. IEEE, 1994, pp. 1–357.
- [49] T. Brants, A. Popat, P. Xu, F. J. Och, and J. Dean, “Large language models in machine translation,” in *Proceedings of the 2007 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning (EMNLP-CoNLL)*, 2007, pp. 858–867.
- [50] T. Mikolov, M. Karafiát, L. Burget, J. Cernocký, and S. Khudanpur, “Recurrent neural network based language model,” in *Interspeech*, vol. 2, no. 3. Makuhari, 2010, pp. 1045–1048.
- [51] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” *Advances in neural information processing systems*, vol. 30, 2017.
- [52] J. Howard and S. Ruder, “Universal language model fine-tuning for text classification,” *arXiv preprint arXiv:1801.06146*, 2018.
- [53] Y. Liu and M. Lapata, “Text summarization with pretrained encoders,” *arXiv preprint arXiv:1908.08345*, 2019.
- [54] J. Yang, M. Wang, H. Zhou, C. Zhao, W. Zhang, Y. Yu, and L. Li, “Towards making the most of bert in neural machine translation,” in *Proceedings of the AAAI conference on artificial intelligence*, vol. 34, no. 05, 2020, pp. 9378–9385.
- [55] J. D. M.-W. C. Kenton and L. K. Toutanova, “Bert: Pre-training of deep bidirectional transformers for language understanding,” in *Proceedings of NAACL-HLT*, 2019, pp. 4171–4186.
- [56] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, “Roberta: A robustly optimized bert pretraining approach,” *arXiv preprint arXiv:1907.11692*, 2019.
- [57] D. Guo, S. Ren, S. Lu, Z. Feng, D. Tang, S. Liu, L. Zhou, N. Duan, A. Svyatkovskiy, S. Fu *et al.*, “Graphcodebert: Pre-training code representations with data flow,” *arXiv preprint arXiv:2009.08366*, 2020.
- [58] M. Lewis, Y. Liu, N. Goyal, M. Ghazvininejad, A. Mohamed, O. Levy, V. Stoyanov, and L. Zettlemoyer, “Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension,” in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 2020, pp. 7871–7880.
- [59] Y. Wang, W. Wang, S. Joty, and S. C. Hoi, “Codet5: Identifier-aware unified pre-trained encoder-decoder models for code understanding and generation,” *arXiv preprint arXiv:2109.00859*, 2021.
- [60] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, I. Sutskever *et al.*, “Language models are unsupervised multitask learners,” *OpenAI blog*, vol. 1, no. 8, p. 9, 2019.
- [61] Y. Dong, X. Jiang, Z. Jin, and G. Li, “Self-collaboration code generation via chatgpt,” *arXiv preprint arXiv:2304.07590*, 2023.
- [62] C. S. Xia and L. Zhang, “Keep the conversation going: Fixing 162 out of 337 bugs for \$0.42 each using chatgpt,” *arXiv preprint arXiv:2304.00385*, 2023.
- [63] K. R. Varshney and H. Alemzadeh, “On the safety of machine learning: Cyber-physical systems, decision sciences, and data products,” *Big data*, vol. 5, no. 3, pp. 246–255, 2017.
- [64] M. Barreno, B. Nelson, A. D. Joseph, and J. D. Tygar, “The security of machine learning,” *Machine Learning*, vol. 81, pp. 121–148, 2010.
- [65] S. Lo Piano, “Ethical principles in machine learning and artificial intelligence: cases from the field and possible ways forward,” *Humanities and Social Sciences Communications*, vol. 7, no. 1, pp. 1–7, 2020.
- [66] J. Yang, K. Zhou, Y. Li, and Z. Liu, “Generalized out-of-distribution detection: A survey,” *arXiv preprint arXiv:2110.11334*, 2021.
- [67] S. Rabanser, S. Günnemann, and Z. Lipton, “Failing loudly: An empirical study of methods for detecting dataset shift,” *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [68] C. Ning and F. You, “Optimization under uncertainty in the era of big data and deep learning: When machine learning meets mathematical programming,” *Computers & Chemical Engineering*, vol. 125, pp. 434–448, 2019.
- [69] A. Ashukha, A. Lyzhov, D. Molchanov, and D. Vetrov, “Pitfalls of in-domain uncertainty estimation and ensembling in deep learning,” *arXiv preprint arXiv:2002.06470*, 2020.
- [70] Y.-C. Hsu, Y. Shen, H. Jin, and Z. Kira, “Generalized odin: Detecting out-of-distribution image without learning from out-of-distribution data,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020, pp. 10 951–10 960.
- [71] A. Kendall and Y. Gal, “What uncertainties do we need in bayesian deep learning for computer vision?” *Advances in neural information processing systems*, vol. 30, 2017.
- [72] J. Gawlikowski, C. R. N. Tassi, M. Ali, J. Lee, M. Humt, J. Feng, A. Kruspe, R. Triebel, P. Jung, R. Roscher *et al.*, “A survey of uncertainty in deep neural networks,” *arXiv preprint arXiv:2107.03342*, 2021.
- [73] P. Oberdiek, M. Rottmann, and H. Gottschalk, “Classification uncertainty of deep neural networks based on gradient information,” in *Artificial Neural Networks in Pattern Recognition: 8th IAPR TC3 Workshop, ANNPR 2018, Siena, Italy, September 19–21, 2018, Proceedings 8*. Springer, 2018, pp. 113–125.
- [74] B. Lakshminarayanan, A. Pritzel, and C. Blundell, “Simple and scalable predictive uncertainty estimation using deep ensembles,” *Advances in neural information processing systems*, vol. 30, 2017.

- [75] D. Barber and C. M. Bishop, "Ensemble learning in bayesian neural networks," *Nato ASI Series F Computer and Systems Sciences*, vol. 168, pp. 215–238, 1998.
- [76] A. Lyzhov, Y. Molchanova, A. Ashukha, D. Molchanov, and D. Vetrov, "Greedy policy search: A simple baseline for learnable test-time augmentation," in *Conference on Uncertainty in Artificial Intelligence*. PMLR, 2020, pp. 1308–1317.
- [77] Y. Gal and Z. Ghahramani, "Dropout as a bayesian approximation: Representing model uncertainty in deep learning," in *international conference on machine learning*. PMLR, 2016, pp. 1050–1059.
- [78] G. Wang, W. Li, M. Aertsen, J. Deprest, S. Ourselin, and T. Vercauteren, "Aleatoric uncertainty estimation with test-time augmentation for medical image segmentation with convolutional neural networks," *Neurocomputing*, vol. 338, pp. 34–45, 2019.
- [79] X. Huang, W. Ruan, W. Huang, G. Jin, Y. Dong, C. Wu, S. Bensalem, R. Mu, Y. Qi, X. Zhao *et al.*, "A survey of safety and trustworthiness of large language models through the lens of verification and validation," *arXiv preprint arXiv:2305.11391*, 2023.
- [80] X. Chen, J. Ye, C. Zu, N. Xu, R. Zheng, M. Peng, J. Zhou, T. Gui, Q. Zhang, and X. Huang, "How robust is gpt-3.5 to predecessors? a comprehensive study on language understanding tasks," *arXiv preprint arXiv:2303.00293*, 2023.
- [81] M. Jang and T. Lukasiewicz, "Consistency analysis of chatgpt," *arXiv preprint arXiv:2303.06273*, 2023.
- [82] Y. Xiao and W. Y. Wang, "On hallucination and predictive uncertainty in conditional language generation," in *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, 2021, pp. 2734–2744.
- [83] A. Malinin and M. Gales, "Uncertainty estimation in autoregressive structured prediction," in *International Conference on Learning Representations*.
- [84] L. Kuhn, Y. Gal, and S. Farquhar, "Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation," *arXiv preprint arXiv:2302.09664*, 2023.
- [85] J. Alammari, "Ecco: An open source library for the explainability of transformer language models," in *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing: System Demonstrations*. Association for Computational Linguistics, 2021.
- [86] A. Galassi, M. Lippi, and P. Torrioni, "Attention in natural language processing," *IEEE transactions on neural networks and learning systems*, vol. 32, no. 10, pp. 4291–4308, 2020.
- [87] J. Vig and Y. Belinkov, "Analyzing the structure of attention in a transformer language model," in *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, 2019, pp. 63–76.
- [88] "Gpt 3.5," <https://platform.openai.com/docs/models/gpt-3-5>, 2023.
- [89] D. Hendrycks and K. Gimpel, "A baseline for detecting misclassified and out-of-distribution examples in neural networks," in *5th International Conference on Learning Representations, ICLR 2017, Toulon, France, April 24–26, 2017, Conference Track Proceedings*. OpenReview.net, 2017. [Online]. Available: <https://openreview.net/forum?id=Hkg4TI9xl>
- [90] U. Arora, W. Huang, and H. He, "Types of out-of-distribution texts and how to detect them," in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, 2021, pp. 10 687–10 701.
- [91] Y. Gal, "Uncertainty in deep learning." University of Cambridge., 2016.
- [92] Z. Wang, Y. Huang, L. Ma, H. Yokoyama, S. Tokumoto, and K. Munakata, "An exploratory study of ai system risk assessment from the lens of data distribution and uncertainty," *arXiv preprint arXiv:2212.06828*, 2022.
- [93] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu, "Bleu: a method for automatic evaluation of machine translation," in *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, 2002, pp. 311–318.
- [94] C. Meister, T. Pimentel, G. Wiher, and R. Cotterell, "Locally typical sampling," *Transactions of the Association for Computational Linguistics*, vol. 11, pp. 102–121, 2023.
- [95] Y. Yang, W.-t. Yih, and C. Meek, "WikiQA: A challenge dataset for open-domain question answering," in *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*. Lisbon, Portugal: Association for Computational Linguistics, Sep. 2015, pp. 2013–2018. [Online]. Available: <https://aclanthology.org/D15-1237>
- [96] C.-Y. Lin, "ROUGE: A package for automatic evaluation of summaries," in *Text Summarization Branches Out*. Barcelona, Spain: Association for Computational Linguistics, Jul. 2004, pp. 74–81. [Online]. Available: <https://aclanthology.org/W04-1013>
- [97] S. Ren, D. Guo, S. Lu, L. Zhou, S. Liu, D. Tang, N. Sundaresan, M. Zhou, A. Blanco, and S. Ma, "Codebleu: a method for automatic evaluation of code synthesis," *arXiv preprint arXiv:2009.10297*, 2020.
- [98] J. Gao, Q. Zhou, and R. Qiu, "ELI5-Category: a categorized open-domain qa dataset," 2021.
- [99] R. Nallapati, B. Zhou, C. Gulcehre, B. Xiang *et al.*, "Abstractive text summarization using sequence-to-sequence rnns and beyond," *arXiv preprint arXiv:1602.06023*, 2016.
- [100] O. Bojar, C. Buck, C. Federmann, B. Haddow, P. Koehn, J. Leveling, C. Monz, P. Pecina, M. Post, H. Saint-Amand, R. Soricut, L. Specia, and A. Tamchyna, "Findings of the 2014 workshop on statistical machine translation," in *Proceedings of the Ninth Workshop on Statistical Machine Translation*. Baltimore, Maryland, USA: Association for Computational Linguistics, Jun. 2014, pp. 12–58. [Online]. Available: <https://aclanthology.org/W14-3302>
- [101] M. Chen, J. Tworek, H. Jun, Q. Yuan, H. P. d. O. Pinto, J. Kaplan, H. Edwards, Y. Burda, N. Joseph, G. Brockman *et al.*, "Evaluating large language models trained on code," *arXiv preprint arXiv:2107.03374*, 2021.
- [102] J. Austin, A. Odena, M. Nye, M. Bosma, H. Michalewski, D. Dohan, E. Jiang, C. Cai, M. Terry, Q. Le *et al.*, "Program synthesis with large language models," *arXiv preprint arXiv:2108.07732*, 2021.
- [103] N. Reimers and I. Gurevych, "Sentence-bert: Sentence embeddings using siamese bert-networks," in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 11 2019. [Online]. Available: <https://arxiv.org/abs/1908.10084>
- [104] X. Chen, M. Lin, N. Schärli, and D. Zhou, "Teaching large language models to self-debug," *arXiv preprint arXiv:2304.05128*, 2023.
- [105] K. Cobbe, V. Kosaraju, M. Bavarian, M. Chen, H. Jun, L. Kaiser, M. Plappert, J. Tworek, J. Hilton, R. Nakano *et al.*, "Training verifiers to solve math word problems," *arXiv preprint arXiv:2110.14168*, 2021.
- [106] "all-mpnet-base-v2," <https://huggingface.co/sentence-transformers/all-mpnet-base-v2>, 2023.